



Original handoff Math reproduction — frozen target LoRA

Matched inference throughput



Accepted progress



128 requests; max 2,048 tokens, natural EOS; one training seed and timing run. Not a model-transfer result.