

RELAYSPEC: REUSING FROZEN DRAFTERS ACROSS LANGUAGE MODEL SIZES

Anonymous authors

Paper under double-blind review

ABSTRACT

A speculative drafter trained on one language model expects that model’s features. Reusing it with a new target therefore requires a compatible input. RelaySpec learns a linear map from the new target’s hidden states to the context consumed by the frozen drafter. The source model supplies fitting features, then leaves the deployed conditioning path. We fit the map on 4,096 problem-and-solution records and retarget Qwen3-4B DFlash and EAGLE-3 drafters to Qwen3-8B and Qwen3-14B. On MATH-500, measured end-to-end throughput is 2.35–5.11 times plain autoregressive decoding. The relay retains 89.1–90.3% of native target-specific drafter throughput in paired comparisons. Math accuracy differs from AR by -0.2 to $+0.8$ percentage points. Fitting-budget, normalization and block-size studies relate accepted progress to throughput. Separate EAGLE-3 measurements show 7.63–7.80 GiB lower peak allocated memory than source reuse.

1 INTRODUCTION

A developer may have a DFlash drafter trained for Qwen3-4B but want to serve Qwen3-8B or Qwen3-14B. The new target can check its proposals, but supplies different hidden states. This creates a deployment choice: train a new target-specific drafter, or reuse the existing drafter while paying for its source model to provide compatible features. RelaySpec addresses deployments that seek to reuse this training investment with limited additional fitting work. It learns a replacement for the source conditioning path, allowing the new target to supply the drafter’s input. The practical goal is faster generation than plain AR while retaining much of a native target-specific drafter’s throughput.

Speculative decoding accelerates autoregressive (AR) generation by checking several proposed tokens in one target pass (Leviathan et al., 2023). DFlash and EAGLE-3 improve proposals using features computed inside a language model (Chen et al., 2026; Li et al., 2025). We call the model that provided these features during drafter training the *source*. The *target* is the model whose token choices govern deployment. Running the source alongside a new target supplies compatible features, but also retains its transformer computation and memory.

RelaySpec runs the source and new target on shared text during fitting. It learns a linear map from the target’s hidden states to the context that the source supplies to the drafter. At inference, the target already computes these states when checking proposals. Applying the map supplies the next drafting input without another source-transformer pass. The existing target verifier still decides which tokens to commit.

Only the map is trained. The target, drafter, inherited embedding and vocabulary projection stay fixed. We apply this procedure to DFlash and EAGLE-3, matching each released decoder’s input normalization. We measure speed relative to AR (Figure 1), native-drafter throughput retained, and additional fitting data and work. Source reuse isolates source-removal savings. Ablations test how fitting data, normalization and draft length affect accepted progress and runtime.




Figure 1: MATH-500 throughput relative to plain AR, measured within each paired run. RelaySpec reuses the 4B drafter. The native control uses a drafter released for the evaluated target. Whiskers show RelaySpec’s 95% paired intervals. Native controls show the three evaluated releases. Each runtime uses one request per GPU.

2 RELATED WORK AND POSITIONING

Reusing and adapting drafters. Table 1 compares feature-conditioned drafting with methods that reuse pretrained drafters or adapt them across targets.

Table 1: Relevant drafting and adaptation methods. The table compares procedures, not runtime rankings. Paired speed measurements appear in Table 4.

Approach	Learned operation
DFlash / EAGLE-3	Train a drafter to propose tokens using its target’s hidden features (Chen et al., 2026; Li et al., 2025)
PARD (An et al., 2026)	Adapt an autoregressive model for parallel drafting, independently of a particular target
SD ² (Berdoz et al., 2026)	Learn target-feature steering injected into drafter MLPs, with drafter fine-tuning in the main configuration
OmniDraft (Ramakrishnan et al., 2025)	Adapt a drafter online through hybrid distillation, using an n-gram cache across vocabularies
RelaySpec	Fit new-target features to the existing source-conditioned input of a frozen drafter

DFlash and EAGLE-3 provide the source-trained drafters and native runtime controls in our study. PARD builds a target-independent parallel drafter, whereas RelaySpec fits a target-specific interface to an existing dependent drafter. SD² also uses target features and tests frozen-drafter variants. Its steering objective aligns token distributions. RelaySpec instead regresses the source context already consumed by the drafter, replacing its source-transformer dependency. OmniDraft addresses online adaptation and vocabulary mismatch. RelaySpec fits offline and uses a shared vocabulary, while heterogeneous verification addresses differing vocabularies (Timor et al., 2025a).

Core contributions of RelaySpec. RelaySpec targets a specific upgrade: deploy a larger target while retaining a released feature-conditioned drafter and a small additional fitting budget. Table 2 separates this requirement from general drafter reuse. Our procedure fits the source context on shared text, replaces source-transformer conditioning at inference, and matches the released DFlash and EAGLE-3 input normalization. Only the interface is trained. The target verifier and drafter stay fixed.

Table 2: Capabilities for reusing a released feature-conditioned drafter with a new target. A check denotes an explicit capability of the described procedure. A dash denotes a requirement it does not address, not an impossibility. This is a method comparison, not a speed ranking.

	Reuse existing feature drafter	Keep drafter weights fixed	Keep target weights fixed	No source transformer
Native DFlash / EAGLE-3	–	–	✓	✓
Source reuse (control)	✓	✓	✓	–
PARD	–	–	✓	✓
SD ²	–	✓ [†]	✓	✓
OmniDraft	–	–	✓	✓
RelaySpec	✓	✓	✓	✓

We compare fitting or adaptation, not freezing after training. Native methods train a drafter for the new target.

[†] SD² tests frozen-drafter variants, although its main configuration fine-tunes the drafter. Methods with independent drafters require no source transformer. References and operations are given in Table 1.

The deployment benefit is to reuse a trained head when upgrading the target within a shared vocabulary. We measure AR speedup, native throughput retained and fitting work across two drafter families and two target sizes. Adapters, low-rank updates and model stitching already adapt frozen networks (Houlsby et al., 2019; Hu et al., 2022; Bansal et al., 2021). Our contribution is replacing the source conditioning of existing speculative decoders. Feature agreement alone does not establish equivalent representations (Smith et al., 2025). Appendix G covers broader drafting and execution methods, including RepSpec and Draft-OPD.

3 PRELIMINARIES

3.1 DRAFTING AND TARGET VERIFICATION

An autoregressive model generates one token at a time. Blockwise and speculative methods propose several tokens, then check them in a target pass (Stern et al., 2018; Xia et al., 2023; Leviathan et al., 2023). In greedy decoding, the verifier commits the longest matching prefix and a target-selected next token. RelaySpec preserves this decoder and changes the features supplied to its drafter.

3.2 THROUGHPUT AND ACCEPTED PROGRESS

Our primary metric is end-to-end output throughput: total generated tokens divided by total request time, including prompt processing. Let T_A , T_R and T_N denote this throughput for plain AR, RelaySpec and the native target-specific drafter. We report

$$\text{AR speedup} = T_R/T_A, \quad \text{native throughput retained} = 100 T_R/T_N \%. \quad (1)$$

A retention value of 90% means the reused drafter produces 90% as many tokens per second as the target’s own drafter. It is not the fraction of speedup above AR. We also report the ratio of summed AR request time to summed relay request time. This differs from a throughput ratio when output lengths differ.

Acceptance explains how speed is obtained. We report the pooled number of tokens advanced per proposal-verification cycle, including the implementation’s target-supplied progress. This is distinct from the fraction of proposed draft tokens accepted. Appendix D gives the position-wise curves and precise counting convention.

Let c_R be the average relay cycle time and a_R its average progress. Its decode time per output token is approximately c_R/a_R . If c_A is AR time per token, then the decode-only approximation is

$$\text{AR-relative decode speed} \approx \frac{a_R c_A}{c_R}. \quad (2)$$

Removing source conditioning reduces cycle cost. Imperfect feature matching may reduce progress. The map is useful when the saved computation outweighs the extra cycles. We measure end-to-end speed directly and use this approximation to interpret the component measurements.

4 METHOD

4.1 FIT THE INPUT CONSUMED BY THE FROZEN DRAFTER

We run the source and new target on the same text and collect hidden states after five selected transformer blocks. Concatenating these states gives the target vector h_t and source vector s_t at position t . The drafter’s frozen projection F converts source features to its input width. We learn a bias-free matrix W that supplies this input from h_t . Figure 2 separates fitting from generation.




Figure 2: Fit once, then reuse the frozen drafter. The source and target process identical text during fitting. Only the map is updated. During generation, target features from committed positions feed the next draft.

The correct input depends on the released drafter. DFlash consumes a normalized projection. Let N_{in} denote fixed root-mean-square normalization of joined target features (Zhang & Sennrich, 2019), which rescales a vector by the square root of its mean squared coordinate. Let N_{out} be the released drafter’s frozen output normalization. EAGLE-3 consumes the raw projection. The teacher context c_t and mapped context $\hat{c}_t$ are

$$\begin{aligned} \text{DFlash: } c_t &= N_{\text{out}}(Fs_t), \quad \hat{c}_t = N_{\text{out}}(WN_{\text{in}}(h_t)), \\ \text{EAGLE-3: } c_t &= Fs_t, \quad \hat{c}_t = Wh_t. \end{aligned} \quad (3)$$

Preserving feature magnitude matters for the raw EAGLE-3 input. DFlash’s output normalization is applied to both teacher and prediction. The normalization ablation tests these choices directly.

We fit W by minimizing the squared context error relative to the teacher’s squared magnitude. For a record with n token positions,

$$\mathcal{L}(W) = \frac{1}{n} \sum_{t=1}^n \frac{\|\hat{c}_t - c_t\|_2^2}{\max(\|c_t\|_2^2, \epsilon)}. \quad (4)$$

Here $\|v\|_2^2$ sums the squared coordinates of v , and ϵ is the smallest positive normal FP32 value. Dividing by teacher magnitude prevents larger context vectors from dominating the objective. We average positions within each record and record losses across workers. DFlash’s normalization makes this objective different from raw linear least squares. Both selected maps are fitted with gradient updates.

Fitting uses 4,096 distinct math problem-and-solution records rendered as conversations, with the resulting text truncated to 192 tokens. The supervision is a source context vector. Solution text remains part of the input. The fitted map has 52.43M weights for an 8B target and 65.54M for a 14B target. Appendix A specifies the layers and dimensions.

4.2 DRAFT WITH THE NEW TARGET’S FEATURES

The target first processes the prompt and retains its attention cache and selected hidden states. The map supplies drafter context. The frozen drafter proposes tokens, and the target checks them in

parallel with a causal mask. The decoder commits the longest matching prefix and the target-selected next token, then crops the cache to the committed prefix. Mapped target states supply the next cycle, as shown in Figure 3.




Figure 3: One generation cycle. RelaySpec changes the drafter’s context. The target retains control over committed tokens and the shared decoder retains responsibility for positions, stopping and cache updates.

The standard greedy equivalence argument requires block verification to return the same token choices as single-token evaluation of each identical prefix (Leviathan et al., 2023). Finite-precision implementations can change those choices. We therefore measure task accuracy and token agreement separately from throughput. The map itself never supplies the scores used for target verification.

5 EXPERIMENTS

5.1 EXPERIMENTAL SETUP

Models and runtime. We retarget the released Qwen3-4B DFlash and DeepSpec EAGLE-3 drafters to Qwen3-8B and Qwen3-14B (Yang et al., 2025; DeepSeek-AI, 2026). DFlash uses block length 16. EAGLE-3 uses the supported chain procedure with seven draft positions. Native controls use released target-specific drafters within the same implementation and proposal setting. The main AR comparisons use four NVIDIA L40S GPUs, with one sequential request per worker. Historical block-size and isolated-memory studies use RTX 6000 Ada GPUs. We compare methods within runs on the same hardware. Serving systems such as vLLM and SGLang also optimize cache use and request execution (Kwon et al., 2023; Zheng et al., 2024). Our timing isolates the stated per-request decoder rather than a concurrent serving workload.

Table 3: Evaluation protocol. Main fitting choices remain fixed across workloads. Full identifiers and runtime versions are in Appendix A.

Component	Setting
Primary baseline	Plain greedy AR with the same target, prompt and runtime
Other controls	Native target-specific drafter and optimized source reuse
Generation	BF16, SDPA attention, thinking disabled, maximum 2,048 new tokens
Relay fitting	4,096 records, 192-token cap, 1,024 AdamW updates (Loshchilov & Hutter, 2019), four records/update
Optimizer	Learning rate 6×10^{-4} , zero weight decay, gradient clipping 1.0
Timing	Synchronized request timer, warmup excluded, rotated method order
Uncertainty	10,000 paired resamples, whole conversations for two-turn chat

Workloads and data use. We evaluate math, code and conversation workloads because speculative speed varies across tasks (Xia et al., 2024; Abramovich et al., 2026). MATH-500 supplies 500 questions (Lightman et al., 2024). The breadth suite contains 128 GSM8K questions, 164 HumanEval tasks, 378 MBPP tasks and 80 two-turn MT-Bench conversations (Hendrycks et al., 2021; Cobbe et al., 2021; Chen et al., 2021; Austin et al., 2021; Zheng et al., 2023). The 32-question development set informed interface and loss choices. Historical evaluations retain this exposure history. Fitting

and evaluation problem texts have zero normalized exact matches. Appendix E quantifies related templates using token overlap. All completed workload cells in each reported suite are included.

Quality and pairing. We score the same math outputs used for the AR timing comparison with the pinned Qwen math scorer. Historical code comparisons use EvalPlus base and expanded tests (Liu et al., 2023). MT-Bench measures two-turn generation time, with each method conditioning its second turn on its own first answer. Its bootstrap unit is the conversation. Other confidence intervals resample matched requests and condition on the fitted checkpoint. Model loading and external task scoring sit outside the request timer.

5.2 ACCELERATION OVER PLAIN AUTOREGRESSIVE DECODING

Table 4 reports absolute end-to-end throughput and AR-relative speed. DFlash RelaySpec reaches 186.49 tokens/s at 8B and 116.88 at 14B, compared with AR at 37.37 and 22.88. EAGLE-3 RelaySpec reaches 84.92 and 59.18 tokens/s, compared with 36.15 and 22.88 for AR. The resulting throughput ratios are 4.99, 5.11, 2.35 and 2.59. Each paired interval is above one.

Table 4: MATH-500 end-to-end tokens/s, including prompt processing. Progress R is mean recorded relay progress per cycle. Ratios and intervals use the AR arm of that same run. Native-drafter comparisons are in Table 6.

Drafter	Target	AR tok/s	Source tok/s	Relay tok/s	Relay / AR [95% CI]	Progress R
DFlash	8B	37.37	121.81	186.49	4.99 [4.88, 5.10]	6.98
DFlash	14B	22.88	92.15	116.88	5.11 [4.98, 5.23]	6.49
EAGLE-3	8B	36.15	66.44	84.92	2.35 [2.32, 2.38]	4.33
EAGLE-3	14B	22.88	53.96	59.18	2.59 [2.55, 2.63]	4.01

DFlash’s parallel proposal step and longer accepted progress yield higher throughput than the sequential EAGLE-3 path in these implementations. This is a comparison of the stated released models and decoding settings. The two integrations use different pinned Transformers versions, so their cross-family difference is not an isolated architectural effect.

5.3 RETAINING TARGET-SPECIFIC DRAFTER PERFORMANCE

For the motivating 4B-to-8B case, RelaySpec reaches 186.49 tokens/s while native DFlash-8B reaches 206.58. Thus the reused 4B drafter retains **90.3%** of native throughput after fitting only the map. EAGLE-3 retains 90.2% at 8B and 89.1% at 14B. Figure 4 places this performance beside the scale of the additional fitting stage.




Figure 4: Left: fraction of native target-specific drafter throughput retained by RelaySpec in paired MATH-500 runs. Right: published original DFlash and EAGLE-3 recipe sizes versus the additional RelaySpec fitting set. Recipe counts provide training-scale context, not a matched training experiment or an audit of each released checkpoint.

Table 5: Original drafter training and additional target adaptation. Record counts describe different training stages and objectives.

Stage	Records	Relative to relay fit	Weights learned
Published DFlash recipe (Chen et al., 2026)	~800,000	~195×	Drafter
Published EAGLE-3 recipe (Li et al., 2025)	~532,000	~130×	Drafter
RelaySpec target adaptation	4,096	1×	Feature map

The DFlash recipe reports about 800K target-generated records. The EAGLE-3 recipe uses about 68K ShareGPT and 464K UltraChat entries. Relay fitting uses 0.51% and 0.77% of these respective record counts. This comparison describes marginal adaptation: the original drafter’s training remains inherited. Record lengths, objectives and hardware differ, so these ratios do not represent equivalent reductions in training compute. The evaluated EAGLE-3 checkpoints come from DeepSpec, while the table’s EAGLE-3 count refers specifically to the published training recipe.

5.4 ACCEPTANCE AND THE SOURCE-REMOVAL MECHANISM

Table 6 pairs retained throughput with accepted progress. RelaySpec advances 6.98 tokens per DFlash-8B cycle versus 7.89 for its native drafter. EAGLE-3 advances 4.33 versus 5.40 at 8B and 4.01 versus 5.28 at 14B. Acceptance and speed retention therefore answer different questions: per-cycle work also changes with drafter size and implementation.

Table 6: Native throughput retained and mean recorded progress per cycle. N and R denote native drafting and RelaySpec. Progress includes the implementation’s target-supplied token.

Drafter	Target	Native tok/s	Relay tok/s	Retained [95% CI]	Progress N	Progress R
DFlash	8B	206.58	186.49	90.3% [89.8, 90.8]	7.89	6.98
EAGLE-3	8B	94.12	84.92	90.2% [89.8, 90.7]	5.40	4.33
EAGLE-3	14B	66.39	59.18	89.1% [88.6, 89.7]	5.28	4.01

Against source reuse, the relay’s throughput ratios are 1.531 and 1.268 for DFlash, and 1.278 and 1.097 for EAGLE-3, at 8B and 14B respectively. This secondary comparison holds the inherited drafter fixed and isolates source replacement. The source’s computation disappears from each cycle, while the target’s retained states already contain useful context. The matched EAGLE-3 normalization result in Section 5.8 supports the importance of supplying that context at its expected scale.

5.5 MATH ACCURACY AND OUTPUT AGREEMENT

Table 7 scores the exact outputs timed in Table 4. Observed AR-to-relay accuracy differences range from -0.2 to $+0.8$ percentage points, and each paired interval includes zero. These measured outcomes support the accuracy comparison without treating token equality as a substitute for task quality.

Table 7: MATH-500 accuracy in percent. Difference is RelaySpec minus AR in percentage points. Token matches count identical full output sequences.

Drafter	Target	AR score	Relay score	Difference [95% CI]	Token matches
DFlash	8B	75.8	76.6	+0.8 [-1.4, +3.0]	117/500
DFlash	14B	79.2	79.2	+0.0 [-2.2, +2.0]	113/500
EAGLE-3	8B	75.8	75.6	-0.2 [-2.4, +2.0]	135/500
EAGLE-3	14B	79.2	79.4	+0.2 [-1.8, +2.2]	106/500

In BF16, AR and block verification can select different tokens on an identical prefix. Traces of eight selected DFlash early-divergence prompts show changes in the leading target scores shared by native drafting, source reuse and RelaySpec. This diagnostic is separate from the timed runs. We report

measured quality and agreement rather than claiming bitwise identity from the verifier’s mathematical contract.

5.6 SPEED ACROSS WORKLOADS

Figure 5 and Table 10 report the completed DFlash AR-paired breadth suite. RelaySpec is 1.85–3.69 times faster than AR across these eight cells. Code gains at 14B remain substantial relative to AR even where source reuse is faster than the relay. The useful comparison depends on the deployment choice: AR measures the acceleration obtained, while source reuse measures the benefit of removing its source.




Figure 5: DFlash RelaySpec throughput relative to AR across workloads. Whiskers are paired 95% intervals. Chat intervals resample whole two-turn conversations. Full absolute values are in Table 10.

5.7 FITTING WORK AND DEPLOYMENT MEMORY

The selected maps have 52.43–65.54M trainable weights. Recorded fitting loops take 85.6–118.5 seconds on four RTX 6000 Ada GPUs with models already loaded. The interval covers feature computation and optimization. Loading, checkpoint writing and deployment initialization are separate setup costs. Appendix F gives each measured fitting loop.

Separate EAGLE-3 processes measure memory after removing source layers. Peak allocated memory falls from 25.56 to 17.76 GiB at 8B and from 37.56 to 29.94 GiB at 14B. These are savings of 7.80 and 7.63 GiB. Source token embeddings and vocabulary projections required by a drafter remain part of deployment. The claim concerns the removed source transformer, not every tensor inherited from the source.

5.8 ABLATION STUDIES

5.8.1 FITTING DATA AND CONTINUED OPTIMIZATION

The completed DFlash-8B budget sweep uses 512, 1,024 and 2,048 distinct records for one pass, followed by a configuration with two passes over 4,096 records. The update counts are 128, 256, 512 and 2,048 with four records per update. Figure 6 shows throughput and progress on the same 128-question evaluation set. This measures a joint increase in data coverage and optimization work.




Figure 6: DFlash-8B fitting-budget sweep on 128 fixed questions. The first three points use one pass over 512, 1,024 and 2,048 distinct records. The last uses 8,192 presentations of 4,096 distinct records. Each point is a separately fitted checkpoint. Lines connect measured budgets.

Throughput rises from 136.78 to 162.04, 173.72 and 187.85 tokens/s as the budget grows. Accepted progress rises with it. More fitting can improve the context supplied to a fixed drafter, reducing the number of cycles required for generation. The last configuration continues fitting over an already used manifest, so its gain reflects additional optimization as well as the larger distinct set. Table 15 reports all configurations and their actual budgets.

5.8.2 DRAFT LENGTH AND INTERFACE CHOICES

The block-size experiment directly compares AR, native DFlash and the relay at lengths 8, 16 and 32 on 64 fixed requests. Block 16 gives the highest relay throughput, 210.15 tokens/s. Block 8 retains more source progress proportionally but advances fewer tokens per cycle. Block 32 adds verification width while reducing accepted progress for the released block-16 drafter. This supports keeping its trained block length instead of maximizing the relative gain over source reuse.

Table 8: DFlash-8B block-size ablation on 64 questions, a 512-token cap and RTX 6000 Ada hardware. AR is measured in each run.

Block	AR tok/s	Native draft tok/s	Relay tok/s	Relay / AR	Progress / cycle
8	44.68	175.12	164.90	3.691	4.98
16	44.67	229.87	210.15	4.704	6.48
32	44.76	207.70	131.36	2.935	4.14

The EAGLE-3 normalization comparison holds parameter count, data and 128-update budget fixed. Raw target features reduce relative context error from 0.413 to 0.266 and increase the source-relative throughput ratio from 1.035 to 1.237. The raw interface consumes feature magnitude, which input normalization removes. This gives a concrete reason for preserving scale in the selected map.

For DFlash, the selected relative-error objective improves throughput by 1.009 and 1.011 times over squared error plus cosine distance on the 32-question comparison at 8B and 14B. The ridge and nonlinear variants in Appendix C provide additional fitting alternatives. Their different objectives and capacities make them comparisons of complete recipes rather than isolated tests of the optimizer or nonlinearity.

6 CONCLUSION

RelaySpec fits an input map to reuse a source-trained drafter with a new target. Using 4,096 additional fitting records, the measured Qwen3 cases achieve 2.35–5.11 times AR throughput and retain 89.1–90.3% of native drafter throughput. These results concern fixed checkpoints, a shared vocabulary and greedy single-request decoding. Confidence intervals condition on the fitted map. The fitting-budget sweep varies data and optimization jointly.

AI USE STATEMENT

Generative AI assisted with implementation, code review, literature search, analysis and manuscript editing. Numerical tables and plots are generated from recorded artifacts with checked method, request and scorer identity. Scientific interpretation and submission decisions remain the authors' responsibility.

REPRODUCIBILITY STATEMENT

The appendix specifies model identifiers, feature layers, fitting settings, measurement boundaries, complete workload results and data checks. The accompanying implementation includes configurations, request records, scorer provenance and scripts that regenerate the reported assets.

REFERENCES

- Talor Abramovich, Maor Ashkenazi, Izzy Putterman, Benjamin Chislett, Tiyasa Mitra, Bitu Darvish Rouhani, Ran Zilberstein, and Yonatan Geifman. SPEED-Bench: A unified and diverse benchmark for speculative decoding. In *43rd International Conference on Machine Learning*, 2026. URL <https://arxiv.org/abs/2604.09557>. Accepted paper.
- Zihao An, Huajun Bai, Ziqiong Liu, Dong Li, and Emad Barsoum. PARD: Accelerating LLM inference with low-cost PARallel draft model adaptation. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=XbOyv7iVGL>.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021. URL <https://arxiv.org/abs/2108.07732>.
- Yamini Bansal, Preetum Nakkiran, and Boaz Barak. Revisiting model stitching to compare neural representations. In *Advances in Neural Information Processing Systems 34*, 2021. URL <https://papers.nips.cc/paper/2021/hash/01ded4259d101feb739b06c399e9cd9c-Abstract.html>.
- Frédéric Berdoz, Peer Rheinboldt, and Roger Wattenhofer. Steering pretrained drafters during speculative decoding. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 30067–30075, 2026. doi: 10.1609/aaai.v40i36.40255. URL <https://ojs.aaai.org/index.php/AAAI/article/view/40255>.
- Tianle Cai, Yuhong Li, Zhengyang Geng, Hongwu Peng, Jason D. Lee, Deming Chen, and Tri Dao. Medusa: Simple LLM inference acceleration framework with multiple decoding heads. In *Proceedings of the 41st International Conference on Machine Learning*, 2024. URL <https://proceedings.mlr.press/v235/cai24b.html>.
- Jian Chen, Yesheng Liang, and Zhijian Liu. DFlash: Block diffusion for flash speculative decoding. In *43rd International Conference on Machine Learning*, 2026. URL <https://arxiv.org/abs/2602.06036>. Accepted paper.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021. URL <https://arxiv.org/abs/2107.03374>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021. URL <https://arxiv.org/abs/2110.14168>.
- DeepSeek-AI. DeepSpec: A full-stack codebase for training and evaluating speculative decoding draft models. Official open-source software, 2026. URL <https://github.com/deepseek-ai/DeepSpec>.

- Mostafa Elhoushi, Akshat Shrivastava, Diana Liskovich, Basil Hosmer, Bram Wasti, Liangzhen Lai, Anas Mahmoud, Bilge Acun, Saurabh Agarwal, Ahmed Roman, Ahmed Aly, Beidi Chen, and Carole-Jean Wu. LayerSkip: Enabling Early Exit Inference and Self-Speculative Decoding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024. URL <https://aclanthology.org/2024.acl-long.681/>.
- Yichao Fu, Peter Bailis, Ion Stoica, and Hao Zhang. Break the Sequential Dependency of LLM Inference Using Lookahead Decoding. In *Proceedings of the 41st International Conference on Machine Learning*, 2024. URL <https://proceedings.mlr.press/v235/fu24a.html>.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems 34, Datasets and Benchmarks Track*, 2021. URL <https://openreview.net/forum?id=7Bywt2mQsCe>.
- Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-Efficient Transfer Learning for NLP. In *Proceedings of the 36th International Conference on Machine Learning*, 2019. URL <https://proceedings.mlr.press/v97/houlsby19a.html>.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations*, 2022. URL <https://www.microsoft.com/en-us/research/publication/lora-low-rank-adaptation-of-large-language-models/>.
- Shijing Hu, Jingyang Li, Xingyu Xie, Zhihui Lu, Kim-Chuan Toh, and Pan Zhou. GRIFFIN: Effective token alignment for faster speculative decoding. In *Advances in Neural Information Processing Systems*, 2025. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/b6e67ae290635d0874c4cb43ba2a2cfb-Abstract-Conference.html.
- Feiye Huo, Jianchao Tan, Jiahao Liu, Zixu Jiang, Jiacheng Li, Jingang Wang, Xunliang Cai, and Shengli Sun. RepSpec: Structural re-parameterized draft model training for speculative decoding. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/d70ea003729b440b89a2f958a5554c1f-Abstract-Conference.html.
- Tanishq Kumar, Tri Dao, and Avner May. Speculative speculative decoding. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/1b96f01343ff10150e6719eb163e1536-Abstract-Conference.html.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient Memory Management for Large Language Model Serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, 2023. URL <https://doi.org/10.1145/3600006.3613165>.
- Haodi Lei, Yafu Li, Haoran Zhang, Shunkai Zhang, Qianjia Cheng, Xiaoye Qu, Ganqu Cui, Bowen Zhou, Ning Ding, Yun Luo, and Yu Cheng. Draft-OPD: On-policy distillation for speculative draft models. *arXiv preprint arXiv:2605.29343*, 2026. URL <https://arxiv.org/abs/2605.29343>.
- Yaniv Leviathan, Matan Kalman, and Yossi Matias. Fast inference from transformers via speculative decoding. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 19274–19286, 2023. URL <https://proceedings.mlr.press/v202/leviathan23a.html>.
- Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE: Speculative Sampling Requires Rethinking Feature Uncertainty. In *Proceedings of the 41st International Conference on Machine Learning*, 2024a. URL <https://proceedings.mlr.press/v235/li24bt.html>.

- Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE-2: Faster Inference of Language Models with Dynamic Draft Trees. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, 2024b. URL <https://aclanthology.org/2024.emnlp-main.422/>.
- Yuhui Li, Fangyun Wei, Chao Zhang, and Hongyang Zhang. EAGLE-3: Scaling up inference acceleration of large language models via training-time test. In *Advances in Neural Information Processing Systems* 38, 2025. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/c7b5a35ea98b62512a869c19ea7b03cb-Abstract-Conference.html.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International Conference on Learning Representations*, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/hash/aca97732e30bcf1303bc22ac3924fd16-Abstract-Conference.html.
- Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation. In *Advances in Neural Information Processing Systems* 36, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/43e9d647ccd3e4b7b5baab53f0368686-Abstract-Conference.html.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019. URL <https://arxiv.org/abs/1711.05101>.
- Xupeng Miao, Gabriele Oliaro, Zhihao Zhang, Xinhao Cheng, Zeyu Wang, Zhengxin Zhang, Rae Ying Yee Wong, Alan Zhu, Lijie Yang, Xiaoxiang Shi, Chunan Shi, Zhuoming Chen, Daiyaan Arfeen, Reyna Abhyankar, and Zhihao Jia. SpecInfer: Accelerating Large Language Model Serving with Tree-based Speculative Inference and Verification. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 3*, 2024. URL <https://arxiv.org/abs/2305.09781>.
- Ramchalam Kinattinkara Ramakrishnan, Zhaocong Yuan, Shaojie Zhuo, Chen Feng, Yicheng Lin, Chenzheng Su, and Xiaopeng Zhang. OmniDraft: A cross-vocabulary, online adaptive drafter for on-device speculative decoding. In *Advances in Neural Information Processing Systems*, 2025. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/3c2fe1417eed1c6ff9acf169617981ea-Abstract-Conference.html.
- Damian Smith, Harvey Mannerling, and Antonia Marcu. Functional alignment can mislead: Examining model stitching. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 55972–55998, 2025. URL <https://proceedings.mlr.press/v267/smith25a.html>.
- Mitchell Stern, Noam Shazeer, and Jakob Uszkoreit. Blockwise Parallel Decoding for Deep Autoregressive Models. In *Advances in Neural Information Processing Systems* 31, 2018. URL <https://proceedings.neurips.cc/paper/2018/hash/c4127b9194fe8562c64dc0f5bf2c93bc-Abstract.html>.
- Nadav Timor, Jonathan Mamou, Daniel Korat, Moshe Berchansky, Gaurav Jain, Oren Pereg, Moshe Wasserblat, and David Harel. Accelerating LLM inference with lossless speculative decoding algorithms for heterogeneous vocabularies. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025a. URL <https://proceedings.mlr.press/v267/timor25a.html>.
- Nadav Timor, Jonathan Mamou, Daniel Korat, Moshe Berchansky, Oren Pereg, Moshe Wasserblat, Tomer Galanti, Michal Gordon-Kiwkowitz, and David Harel. Distributed speculative inference: Speculation parallelism for provably faster lossless language model inference. In *International Conference on Learning Representations*, 2025b. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/b36554b97da741b1c48c9de05c73993e-Abstract-Conference.html.

- Jikai Wang, Yi Su, Juntao Li, Qingrong Xia, Zi Ye, Xinyu Duan, Zhefeng Wang, and Min Zhang. OPT-Tree: Speculative Decoding with Adaptive Draft Tree Structure. *Transactions of the Association for Computational Linguistics*, 2025. URL <https://aclanthology.org/2025.tacl-1.8/>.
- Heming Xia, Tao Ge, Peiyi Wang, Si-Qing Chen, Furu Wei, and Zhifang Sui. Speculative Decoding: Exploiting Speculative Execution for Accelerating Seq2seq Generation. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023. URL <https://aclanthology.org/2023.findings-emnlp.257/>.
- Heming Xia, Zhe Yang, Qingxiu Dong, Peiyi Wang, Yongqi Li, Tao Ge, Tianyu Liu, Wenjie Li, and Zhifang Sui. Unlocking Efficiency in Large Language Model Inference: A Comprehensive Survey of Speculative Decoding. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024. URL <https://aclanthology.org/2024.findings-acl.456/>.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025. URL <https://arxiv.org/abs/2505.09388>.
- Biao Zhang and Rico Sennrich. Root Mean Square Layer Normalization. In *Advances in Neural Information Processing Systems 32*, 2019. URL <https://proceedings.neurips.cc/paper/2019/hash/1e8a19426224ca89e83cef47f1e7f53b-Abstract.html>.
- Jun Zhang, Jue Wang, Huan Li, Lidan Shou, Ke Chen, Gang Chen, and Sharad Mehrotra. Draft & Verify: Lossless Large Language Model Acceleration via Self-Speculative Decoding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024. URL <https://aclanthology.org/2024.acl-long.607/>.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information Processing Systems 36, Datasets and Benchmarks Track*, 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/91f18a1287b398d378ef22505bf41832-Abstract-Datasets_and_Benchmarks.html.
- Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng. SGLang: Efficient Execution of Structured Language Model Programs. In *Advances in Neural Information Processing Systems 37*, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31balc9ac630f15dec8-Abstract-Conference.html.
- Xiandong Zou, Jianshu Li, Jing Huang, and Pan Zhou. Variational speculative decoding: Rethinking draft training from token likelihood to sequence acceptance. In *Proceedings of the 43rd International Conference on Machine Learning*, 2026. URL https://ink.library.smu.edu.sg/sis_research/11190/.

APPENDIX

The appendix gives implementation details (Appendix A), complete workload results (Appendix B), fitting studies (Appendix C), acceptance measurements (Appendix D), data checks (Appendix E) and fitting-time and memory measurements (Appendix F).

A IMPLEMENTATION AND MEASUREMENT DETAILS

Released models. The source is Qwen/Qwen3-4B. Targets are Qwen/Qwen3-8B and Qwen/Qwen3-14B. The source DFlash is z-lab/Qwen3-4B-DFlash-b16, and its native 8B control is z-lab/Qwen3-8B-DFlash-b16. EAGLE-3 uses DeepSpec’s eagle3_qwen3_4b_ttt7, eagle3_qwen3_8b_ttt7 and eagle3_qwen3_14b_ttt7 releases. All model and upstream revisions are fixed in the accompanying configurations.

Features and dimensions. Layer indices count transformer blocks from zero. Source and Qwen3-8B states come from blocks [1, 9, 17, 25, 33]. Qwen3-14B states come from [1, 10, 19, 28, 37]. Joining five 4,096-wide states gives a 20,480-wide 8B input. Joining five 5,120-wide states gives a 25,600-wide 14B input. Both drafter interfaces have width 2,560. The matrix dimensions are therefore $2,560 \times 20,480$ and $2,560 \times 25,600$. The DFlash input normalization has no trainable gain or bias. The released output normalization remains frozen.

Training. Four workers average record losses at each update. Seed 1729 fixes the selected fit. AdamW uses learning rate 0.0006, zero weight decay and gradient clipping at 1.0. Computation uses BF16, with the feature loss computed in FP32. Each of 1,024 updates processes four records, each capped at 192 rendered tokens. The same map remains fixed across tasks.

Runtime and timing. The main AR runs use Python 3.12.13 and PyTorch 2.9.1 with CUDA 12.8. DFlash uses Transformers 5.3.0. EAGLE-3 uses the pinned 5.10.2 overlay. Each worker owns a full target replica on one L40S GPU. Four workers process independent requests rather than splitting one model over GPUs. One warmup per method precedes measured requests. Method order rotates across requests. The synchronized timer includes prompt processing, drafting, verification, feature updates and runtime bookkeeping. The historical block and memory experiments use RTX 6000 Ada GPUs.

Stopping and conversation history. Generation stops at the first end token or the output cap. Capped requests remain in both timing and quality aggregates. Each MT-Bench method uses its own first-turn output as second-turn context. Two turns form one resampling unit. The breadth run has 750 initial records and 830 timed requests after expanding 80 conversations to two turns.

Aggregation. Each method’s throughput is the sum of its output-token counts divided by the sum of its request times. Each of 10,000 resamples recomputes this ratio from matched requests, using bootstrap seed 1729. Confidence bounds are the central 95% of resampled estimates. The request-time ratio uses summed times alone. Table 9 reports both measurements and output lengths.

Table 9: Additional MATH-500 measurements. A is AR, S is source reuse and R is RelaySpec. Time ratio divides summed AR time by summed relay time.

Drafter	Target	Mean tokens	AR / R	Time ratio AR / R	Throughput R / S	Progress S / R
DFlash	8B	783.22 / 778.81		5.018	1.531	7.66 / 6.98
DFlash	14B	777.66 / 772.33		5.143	1.268	7.44 / 6.49
EAGLE-3	8B	783.22 / 767.79		2.397	1.278	5.15 / 4.33
EAGLE-3	14B	777.66 / 766.20		2.625	1.097	5.07 / 4.01

B COMPLETE WORKLOAD RESULTS

Table 10: Complete DFlash breadth comparisons with paired AR. Ratios use end-to-end throughput. Every planned task and target cell is included.

Target	Task	Requests	AR tok/s	Relay tok/s	Relay / AR [95% CI]	Relay / source
8B	GSM8K	128	37.60	138.89	3.69 [3.59, 3.81]	1.456
8B	HumanEval	164	37.55	120.87	3.22 [3.14, 3.30]	1.244
8B	MBPP	378	37.49	129.35	3.45 [3.36, 3.54]	1.320
8B	MT-Bench	160	37.47	70.69	1.89 [1.76, 2.05]	1.387
14B	GSM8K	128	26.80	96.98	3.62 [3.52, 3.72]	1.104
14B	HumanEval	164	26.73	77.54	2.90 [2.83, 2.98]	0.890
14B	MBPP	378	26.72	83.82	3.14 [3.07, 3.21]	0.956
14B	MT-Bench	160	26.55	49.26	1.85 [1.73, 2.01]	1.086

The earlier fixed-map source comparisons in Table 11 provide all four workloads for both drafter families on RTX 6000 Ada. They are a separate paired experiment with their own denominator. The EAGLE-3 14B code cells show the reduced margin when lower accepted progress outweighs source-removal savings. This operating range complements the AR acceleration measured in the main comparison.

Table 11: Historical paired source-reuse breadth measurements on RTX 6000 Ada. Ratios compare relay with source throughput within these runs. They are not normalized by an AR arm from another experiment.

Drafter	Target	Task	Requests	Source tok/s	Relay tok/s	Ratio [95% interval]
DFlash	8B	GSM8K	128	116.87	166.39	1.4237 [1.4040, 1.4445]
DFlash	8B	HumanEval	164	118.84	145.00	1.2201 [1.2017, 1.2383]
DFlash	8B	MBPP	378	117.77	151.55	1.2868 [1.2732, 1.3001]
DFlash	8B	MT-Bench	160	60.85	81.62	1.3412 [1.3183, 1.3639]
DFlash	14B	GSM8K	128	88.04	97.07	1.1025 [1.0836, 1.1214]
DFlash	14B	HumanEval	164	87.34	77.61	0.8886 [0.8741, 0.9038]
DFlash	14B	MBPP	378	88.02	83.98	0.9541 [0.9416, 0.9663]
DFlash	14B	MT-Bench	160	45.32	49.30	1.0877 [1.0653, 1.1083]
EAGLE-3	8B	GSM8K	128	85.99	108.45	1.2613 [1.2437, 1.2789]
EAGLE-3	8B	HumanEval	164	71.75	81.10	1.1304 [1.1172, 1.1434]
EAGLE-3	8B	MBPP	378	69.85	82.10	1.1753 [1.1656, 1.1850]
EAGLE-3	8B	MT-Bench	160	42.27	49.45	1.1698 [1.1501, 1.1906]
EAGLE-3	14B	GSM8K	128	67.81	73.41	1.0826 [1.0675, 1.0979]
EAGLE-3	14B	HumanEval	164	55.38	50.58	0.9132 [0.9011, 0.9255]
EAGLE-3	14B	MBPP	378	54.55	50.49	0.9256 [0.9151, 0.9361]
EAGLE-3	14B	MT-Bench	160	33.83	34.57	1.0221 [0.9999, 1.0449]

Table 12: Task scores for the historical fixed-map breadth runs. Source reuse and RelaySpec share each reported score. Code columns are single-program pass percentages under base and expanded EvalPlus tests.

Drafter	Target	GSM8K	HumanEval	HumanEval+	MBPP	MBPP+
DFlash	8B	93.75	85.98	80.49	84.13	73.02
DFlash	14B	94.53	89.02	83.54	88.36	75.13
EAGLE-3	8B	92.97	87.20	82.32	83.07	71.96
EAGLE-3	14B	94.53	90.24	85.98	88.89	74.87

Across these historical math/code comparisons, all 4,680 paired task outcomes agree between source reuse and RelaySpec. This count is four model/drafter pairs times $500 + 128 + 164 + 378 = 1,170$ problems. Expanded code tests evaluate the same programs and do not add independent examples.

For the current AR-paired DFlash GSM8K runs, the scorer gives 120/128 correct for all arms at 8B and 121/128 at 14B. Code scores above refer specifically to the historical outputs, preserving scorer/run identity.

AR-paired EAGLE-3 code quality. We additionally score the saved programs from the EAGLE-3 AR-paired breadth runs with official EvalPlus base and expanded tests. Table 13 compares one program per task against AR from the same run. At 8B, RelaySpec passes four fewer MBPP tasks under both suites. At 14B, it passes two more. HumanEval expanded-test totals match AR at both target sizes, although individual task outcomes differ. RelaySpec and source reuse have identical per-task outcomes in all eight cells. The paired intervals cover zero. This does not establish equality or noninferiority. These retrospective comparisons use one fitted map per setting and do not include fitting-seed uncertainty. Conversation quality remains unmeasured.

Table 13: EAGLE-3 AR-paired code quality. Counts are passed tasks. The intervals resample paired tasks. Expanded tests reuse the same programs and do not add independent samples.

Target/task	Suite	AR	Relay	Δ (pp)	Paired 95% interval
8B HumanEval	base	143/164	141/164	-1.22	$[-3.66, +1.22]$
8B HumanEval	plus	133/164	133/164	+0.00	$[-3.05, +3.05]$
8B MBPP	base	318/378	314/378	-1.06	$[-2.91, +0.79]$
8B MBPP	plus	276/378	272/378	-1.06	$[-2.65, +0.53]$
14B HumanEval	base	143/164	143/164	+0.00	$[-2.44, +2.44]$
14B HumanEval	plus	134/164	134/164	+0.00	$[-2.44, +2.44]$
14B MBPP	base	338/378	340/378	+0.53	$[-1.32, +2.38]$
14B MBPP	plus	283/378	285/378	+0.53	$[-1.06, +2.12]$

C FITTING AND INTERFACE STUDIES

Distinct records at fixed fitting work. A separate controlled DFlash-8B study holds fitting at 1,024 updates and four records per update while varying distinct records from 64 to 4,096. Smaller sets repeat in the same ordered schedule. Table 14 reports end-to-end throughput on the same 128 development-exposed MATH questions. Of the measured settings, 512 is the smallest within 5% of the best observed mapper throughput. The 256-record setting retains 91.4% of that observed best. This is a descriptive, single-seed boundary, not a minimum across settings or a confirmed optimum. The 512–4,096 settings form an unresolved plateau. Selecting the reference from the same evaluations does not provide independent confirmation. Every relay scores 105/128, versus AR’s 104/128.

Table 14: Distinct-data study at 4,096 presentations for every row. Native retention uses each run’s paired native-drafter control. Best retention uses the observed best relay across these settings. All results condition on one fitting seed.

Records	Tokens/s	$\times$ AR	Native retained	Best retained	Correct
64	134.39	3.55	65.7%	72.8%	105/128
128	154.56	4.10	75.5%	83.7%	105/128
256	168.76	4.51	82.6%	91.4%	105/128
512	181.91	4.84	88.9%	98.5%	105/128
1,024	184.55	4.92	90.3%	99.9%	105/128
2,048	183.81	4.88	89.8%	99.5%	105/128
4,096	184.64	4.90	90.2%	100.0%	105/128

Continuous optimization. We also save checkpoints from one uninterrupted 4,096-record fit at 128, 256, 512, 1,024 and 2,048 updates. They have seen respectively 512, 1,024, 2,048, 4,096 and 4,096 distinct records. The final point therefore measures a second pass, not 8,192 distinct examples. End-to-end AR-relative throughput increases from 3.65 to 5.03 across this trajectory. All five checkpoints score 105/128. Unlike the earlier separately fitted budget configurations, these checkpoints share one optimization trajectory. Figure 7 separates this experiment from the fixed-work data study.




Figure 7: Controlled data and optimization studies on 128 fixed development requests. Error bars are paired 95% request-bootstrap intervals within each run. They exclude fitting-seed variation and selection uncertainty. Connecting lines do not imply unmeasured results.

Table 15: Completed DFlash-8B fitting studies on the same 128-question set. Records count distinct fitting examples. Two passes over 4,096 records give 8,192 presentations. Source reuse is paired within each run.

Fitting choice	Records	Updates	Relay tok/s	Relay / source [95% CI]	Score
512 records	512	128	136.78	1.150 [1.118, 1.181]	82.03
1,024 records	1,024	256	162.04	1.372 [1.347, 1.397]	82.03
2,048 records	2,048	512	173.72	1.477 [1.458, 1.495]	82.03
4,096 records, two passes	4,096	2,048	187.85	1.573 [1.558, 1.588]	82.03
Ridge surrogate	4,096	<i>solve</i>	115.18	0.955 [0.927, 0.982]	82.03
Nonlinear width 512	4,096	1,024	104.49	0.872 [0.841, 0.903]	82.03

The data-budget configurations use the same normalized linear interface and relative-error objective. The nonlinear map uses an intermediate width of 512 and about 11.8M parameters, compared with 52.4M for the dense 8B map. Its lower throughput is consistent with a less effective fitted interface, but the comparison also changes capacity. It does not isolate a disadvantage of nonlinear functions. The ridge configuration solves a raw linear surrogate with ridge coefficient 1.0. Since the deployed DFlash interface includes output normalization, this is a different objective from Equation 4.

Matched capacity on a separate fitting source. We additionally fit 30 DFlash-8B configurations on nested sets of 512 and 2,048 NuminaMath-CoT records,¹ using a fixed 1,024-record fitting validation split. The source-stratified sets include `cn_k12`, `orca_math` and `synthetic_math` examples. Their lexical filter excludes exact normalized problem/solution matches and near matches at five-shingle Jaccard similarity at least 0.6 for texts with at least ten unique shingles. This filter does not establish semantic or template independence. These results use a separate fitting source from the historical MATH study above.

Each configuration receives 8,192 updates with four records per update: 64 passes at $N = 512$ and 16 at $N = 2048$. Frozen features are cached once. Factored linear maps and bias-free two-layer GELU MLPs use widths 64, 128, 256, 512, 1,024, 2,048 and 4,096, giving matched parameter counts $h(20,480 + 2,560)$. The dense reference has 52.43M parameters. The rank of a linear function remains bounded by 2,560 even when its factorization width is larger. All cells share the data order, relative-interface objective, learning rate 6×10^{-4} and seed 1729.

¹AI-MO/NuminaMath-CoT, revision 9d8d210c9f6a, training shard 0 of 5.




Figure 8: Complete primary capacity study at 8,192 updates per cell. Markers compare validation interface error, with one fitting seed. The curves do not establish decoding speed or answer quality.

Table 16: Validation relative interface MSE for the capacity study. Linear and MLP entries at the same width have equal parameter counts. The dense row supplies the single-linear-layer reference.

Width	Params (M)	Linear, $N=512$	MLP, $N=512$	Linear, $N=2048$	MLP, $N=2048$
64	1.47	0.5732	0.5790	0.5692	0.5752
128	2.95	0.4761	0.4871	0.4690	0.4804
256	5.90	0.3760	0.3883	0.3648	0.3767
512	11.80	0.2862	0.3082	0.2685	0.2878
1024	23.59	0.2295	0.2579	0.2040	0.2224
2048	47.19	0.2133	0.2438	0.1847	0.1982
4096	94.37	0.2145	0.2501	0.1845	0.1936
Dense	52.43	0.2101	—	0.1805	—

At this optimization budget, wider factorizations approach the dense reference, while MLPs have higher validation error at every matched width. Increasing MLP width from 2,048 to 4,096 at $N = 512$ lowers diagnostic training error from 0.1586 to 0.1278 but raises validation error from 0.2438 to 0.2501. At $N = 2048$, validation error instead decreases from 0.1982 to 0.1936. This interaction motivates checking regularization and convergence before attributing the ranking to function class.

More epochs on fixed smaller datasets. We continue the width-512 linear and MLP fits to 32,768 total updates, preserving Adam moments, random states and cyclic data position. A separate bounded pilot reproduced uninterrupted fitting bit for bit. Figure 9 shows late-training diagnostics on the same first 256 training records and all 1,024 validation records. The $N = 2048$ MLP improves from 0.2878 to 0.2787 validation error between 8,192 and 32,768 updates. At $N = 512$, linear training error falls from 0.2443 to 0.2350 while validation error rises from 0.2862 to 0.2873. The latter is consistent with mild overfitting on this diagnostic split. These single-seed fitting trajectories require separate decoding and quality evaluation, including replication across targets and workloads.




Figure 9: Late-training trajectories, starting at 2,048 updates. Dashed lines use the fixed training diagnostic subset. Solid lines use fitting validation. The horizontal axis counts updates, while the distinct datasets remain fixed at 512 or 2,048 records.

Explicit regularization on fixed smaller datasets. We complete 42 nonzero-penalty fits for the dense, width-512 factored linear and width-512 GELU MLP connectors at $N \in \{512, 2048\}$. Each fit uses 8,192 updates, the same 1,024 validation records and seed 1729. Explicit L2 coefficients are $10^{-7}, 10^{-6}, 10^{-5}, 10^{-4}$. Separate AdamW arms use decay $10^{-4}, 10^{-3}, 10^{-2}$. No tested L2 coefficient improves its matched zero-penalty endpoint. The largest AdamW improvement is a 0.17% relative decrease in validation error for the width-512 linear connector at $N = 512$. Neither MLP setting improves with either penalty family. Table 17 reports the best validation value within each penalty family, so these minima are development selections. Figure 10 retains all tested coefficients. This grid does not establish a decoding gain or a penalty effect across seeds, and it does not resolve regularization for the wider MLPs.

Table 17: Matched regularization endpoints. Lower relative interface MSE is better. Each nonzero minimum is selected from four L2 or three AdamW coefficients using the same fitting validation split.

N	Mapper	No penalty	Best tested L2	Best tested AdamW
512	Dense	0.21008	0.21138	0.21011
512	Factored linear 512	0.28617	0.28738	0.28570
512	GELU MLP 512	0.30818	0.30827	0.30822
2048	Dense	0.18055	0.18424	0.18055
2048	Factored linear 512	0.26853	0.26970	0.26852
2048	GELU MLP 512	0.28776	0.28864	0.28779




Figure 10: All 42 penalty endpoints relative to their matched zero-penalty validation error. Positive values indicate worse fitting validation. The two curves are separate objective-L2 and AdamW experiments.

Table 18: Matched EAGLE-3 8B normalization comparison on 32 development questions after 128 updates. Throughput ratios use source reuse.

Input to linear map	Relative error	Throughput ratio [95% interval]	Progress retained
Normalized features	0.413	1.035 [1.003, 1.068]	69.0%
Raw features	0.266	1.237 [1.213, 1.263]	82.6%

Table 19: DFlash objective comparison on 32 development questions. Candidates share fitting data, update count, normalization and optimizer. Each ratio uses the source-reuse arm of that candidate’s paired run.

Target	Objective	Throughput ratio [95% interval]	Token matches
8B	Relative error	1.523 [1.488, 1.558]	32/32
8B	Squared error + cosine	1.509 [1.476, 1.545]	32/32
14B	Relative error	1.265 [1.243, 1.289]	32/32
14B	Squared error + cosine	1.251 [1.229, 1.276]	32/32

Relative error weights a record’s feature mismatch by teacher magnitude. Squared error plus 0.1 times cosine distance combines coordinate error with directional agreement. Direct candidate-to-candidate throughput ratios are 1.0085 [1.0005, 1.0174] at 8B and 1.0114 [1.0046, 1.0177] at 14B. These are modest gains with one fewer mixing coefficient.

C.1 EAGLE-3 FITTING REPLICATION

We repeat the fixed smaller-data fitting study with the EAGLE-3 source drafter and the 8B target. The grid contains dense, factored linear and GELU MLP connectors, with widths 128, 512 and 2048 for the latter two. Each primary fit uses 8,192 updates at $N \in \{512, 2048\}$ and the same 1,024 fitting-validation records. Two additional seeds repeat the dense $N = 2048$ setting. This experiment uses the EAGLE input convention without the DFlash input normalization. Error values are interpreted within each drafter family’s conditioning interface.

Table 20: EAGLE-3 fitting validation. End denotes the 8,192-update endpoint. Best denotes the minimum saved fitting-validation error, with its update count. These are development selections for seed 1729, not decoding results. M weights counts millions of trainable mapper parameters.

Mapper	M weights	$N = 512$			$N = 2048$		
		End	Best	Update	End	Best	Update
Dense	52.43	0.24787	0.21666	1024	0.21203	0.21061	2048
Linear 128	2.95	0.41980	0.41980	8192	0.41505	0.41505	8192
Linear 512	11.80	0.28632	0.28505	2048	0.27188	0.27188	8192
Linear 2048	47.19	0.23486	0.22877	2048	0.21132	0.21132	8192
MLP 128	2.95	0.46554	0.46554	8192	0.45645	0.45645	8192
MLP 512	11.80	0.31749	0.31749	8192	0.29577	0.29577	8192
MLP 2048	47.19	0.28154	0.27124	2048	0.23196	0.22683	4096

The dense $N = 512$ mapper reaches 0.21666 at 1,024 updates, then rises to 0.24787 at the final update. Its endpoint is worse than the width-2048 linear mapper, but its best saved validation value is better than that linear mapper’s 0.22877. At $N = 2048$, the dense and width-2048 linear minima are 0.21061 and 0.21132. Thus endpoint rankings alone can reflect different convergence trajectories. At every matched width in this grid, the MLP’s best saved validation error exceeds the linear mapper’s minimum. The three dense $N = 2048$ endpoints range from 0.212034 to 0.212043. This narrow seed spread concerns one fitting setting. It does not establish seed stability for the other architectures or for decoding performance.




Figure 11: EAGLE-3 fitting-validation trajectories at fixed distinct-data counts. Additional updates can worsen validation, most visibly for the dense mapper at $N = 512$. The horizontal axis counts updates rather than new examples. The subsequent checkpoint comparison measures actual decoding.

C.2 FEATURE-VALIDATION MINIMA AND DECODING

We compare each distinct saved feature-validation minimum with its 8,192-update endpoint in two declared decoding campaigns. Each uses 16 paired development questions, greedy decoding and a 256-token cap. Checkpoint selection precedes decoding. Table 21 includes every distinct early-versus-endpoint pair in these campaigns.

Table 21: Decoding throughput at the minimum feature-validation checkpoint relative to the same fit’s endpoint. Intervals resample paired requests 10,000 times. They are descriptive, without multiple-comparison adjustment, and do not measure variation across fitting seeds.

Family	Mapper	N	Learning rate	Selected update	Early / end [95% CI]
DFlash	Linear 4096	512	0.0006	4096	1.009 [0.981, 1.040]
DFlash	MLP 4096	512	0.0006	4096	0.978 [0.961, 0.998]
DFlash	Linear 4096	512	0.0002	4096	1.008 [0.985, 1.029]
DFlash	Linear 4096	512	0.0018	4096	1.003 [0.978, 1.029]
DFlash	MLP 4096	512	0.0002	2048	1.011 [0.986, 1.039]
DFlash	MLP 4096	512	0.0018	4096	0.961 [0.939, 0.983]
EAGLE-3	Dense	512	0.0006	1024	0.950 [0.924, 0.976]
EAGLE-3	Linear 512	512	0.0006	2048	0.969 [0.948, 0.988]
EAGLE-3	Linear 2048	512	0.0006	2048	0.964 [0.922, 1.003]
EAGLE-3	MLP 2048	512	0.0006	2048	0.968 [0.942, 0.996]
EAGLE-3	Dense	2048	0.0006	2048	0.980 [0.969, 0.992]
EAGLE-3	MLP 2048	2048	0.0006	4096	1.019 [0.995, 1.046]

Lower feature-validation loss does not consistently predict faster decoding. For EAGLE-3’s dense $N = 512$ fit, the selected 1,024-update checkpoint retains 94.99% of endpoint throughput, with interval [92.39, 97.59]%. For the DFlash width-4096 MLP at $N = 512$ and learning rate 0.0006, the selected 4,096-update checkpoint retains 97.77%, with interval [96.09, 99.76]%. Other intervals cross equal throughput. These results support treating feature regression as a surrogate and measuring decoding when selecting a connector. They do not establish full-answer quality or justify selection on untouched evaluation data.

Table 22: Observed throughput ranges across three fitting seeds in a separate 16-question, 256-token development campaign. Selected uses each fit’s minimum saved feature-validation checkpoint. These ranges describe the three measured fits and runtime variation, not confidence intervals.

Mapper	N	Learning rate	Endpoint tok/s range	Selected tok/s range
Dense	512	0.0006	184.34–184.68	183.36–184.58
Dense	2048	0.0006	186.11–189.22	186.11–189.22
Linear 1024	512	0.0006	179.47–181.23	179.47–181.23
Linear 4096	2048	0.0002	185.36–187.81	185.36–187.81
MLP 4096	2048	0.0002	183.92–184.78	183.92–184.78
MLP 4096	512	0.0002	174.53–177.30	175.57–177.56

All six settings in Table 22 use the same cached features and seeds 1729, 1730 and 1731. The dense $N = 512$ endpoints span 184.34–184.68 tokens/s. The width-4096 $N = 2048$ MLP improves on its $N = 512$ counterpart, while the dense connector remains competitive. These short development runs do not establish the minimum sufficient data count or distinguish small architecture gaps from runtime variation.

C.3 LONGER DEVELOPMENT OUTPUTS FOR SMALL-DATA MAPPERS

We evaluate nine saved DFlash connectors on 128 paired MATH development questions with a 2,048-token cap, greedy decoding and the same source, target and native-drafter references. These exposed questions provide a larger development comparison than the 16-question, 256-token screen. They are separate from any untouched confirmation set. The continued width-512 fits reuse the same 2,048 distinct training examples.

Table 23: Capped development outcomes for small-data DFlash fitting. Throughput retention uses the dense $N = 2048$, 8,192-update connector. Intervals use 10,000 paired request bootstrap samples. Correct counts scored answers. Cap counts outputs reaching 2,048 tokens, which can include EOS exactly at the limit and are not exact truncation counts.

Method	N	Updates	Tok/s	Dense retained (%) [95% CI]	Correct	Cap
AR	–	–	37.72	19.04 [18.23, 19.90]	104/128	11
Native drafter	–	–	204.57	103.27 [102.35, 104.17]	105/128	9
Source reuse	–	–	120.96	61.06 [60.53, 61.59]	105/128	9
Dense	512	8192	193.68	97.78 [97.09, 98.47]	105/128	9
Dense	2048	8192	198.09	100.00 [100.00, 100.00]	105/128	9
Linear 1024	512	8192	188.19	95.00 [94.14, 95.85]	105/128	9
Linear 4096	2048	8192	196.25	99.07 [98.50, 99.60]	105/128	9
MLP 4096	2048	8192	190.15	95.99 [95.30, 96.70]	105/128	9
Linear 512	2048	8192	164.61	83.10 [81.81, 84.38]	105/128	9
Linear 512	2048	32768	167.29	84.46 [83.15, 85.71]	105/128	9
MLP 512	2048	8192	154.39	77.94 [76.50, 79.35]	105/128	9
MLP 512	2048	32768	157.68	79.60 [78.07, 81.06]	105/128	9

Dense fitting on 512 examples retains 97.78% of the dense $N = 2048$ throughput, with interval [97.09, 98.47]%. This establishes a lower-data performance point in this setting, but does not establish that 512 is minimal. The width-1024 linear mapper at $N = 512$ retains 95.00%, with interval [94.14, 95.85]%, leaving its 95% boundary unresolved. The width-4096 linear mapper is slightly below the dense connector in this larger sample, so the earlier short-screen point ranking should not be interpreted as an established architecture advantage.

All nine connectors score 105 of 128 answers, compared with 104 for AR. Their paired accuracy-difference intervals are $[-0.03125, 0.046875]$. These intervals do not establish a one-percentage-point noninferiority margin. Nine connector outputs and eleven AR outputs reach the token cap. The observed scores therefore concern the configured capped evaluation, not uncapped answer quality. Longer width-512 fitting gives modest throughput increases while remaining below the wider or dense connectors. This comparison does not treat additional updates as additional data.

D ACCEPTANCE AND NUMERICAL AGREEMENT

We pool all recorded proposal cycles rather than averaging each request’s mean. This weights cycles equally. DFlash’s recorded length is one already target-selected position plus the accepted draft prefix. EAGLE-3’s length follows its shared decoder’s progress count. These quantities measure work advanced per cycle, including target-supplied progress. A proposed-token acceptance fraction requires separating those positions and accounting for the number proposed.




Figure 12: Historical 8B source-reuse comparisons. Each curve gives the fraction of recorded cycles reaching each position. Summing the curve values gives mean recorded progress per cycle. Solid and dashed lines show source reuse and RelaySpec, respectively.

The main MATH-500 table reports measured full-sequence agreement with AR. A fresh diagnostic selects eight prompts with early differences and saves target inputs, cache lengths, positions and the two leading logits at every call. The DFlash paths share the same first differing choice across native, source and relay drafting. The observed BF16 logits include ties or changed ordering with gaps up to 0.25 at those positions. Such traces identify numerical score differences in the selected cases. They do not establish that all mismatches have the same cause. The task-quality intervals in Table 7 remain the empirical quality evidence.

E DATA CHECKS AND DEVELOPMENT EXPOSURE

The 4,096-record fitting manifest contains distinct problems with solutions. The loader renders both fields before applying the 192-token cap. Repeating the manifest changes presentations, not distinct records. The 32-question development set informed normalization and objective choices. The historical 468-question complement is computed from saved full-run outputs and retains that exposure history.

Normalized exact matching finds zero overlaps between the 4,096 fitting and 1,250 evaluation problem texts. For a second check, each text is split into unique normalized tokens. Token-set Jaccard overlap is the number of shared tokens divided by the number of tokens in either set. Each evaluation question is compared with its closest fitting record. Table 24 reports the resulting similarities.

Table 24: Closest fitting-to-evaluation token-set overlap. Threshold columns count evaluation records at or above the stated overlap.

Task	Maximum overlap	≥ 0.80	≥ 0.90	≥ 0.95
MATH-500	0.980	47	13	4
GSM8K	0.378	0	0	0
HumanEval	0.290	0	0	0
MBPP	0.333	0	0	0
MT-Bench	0.571	0	0	0

MATH-500 includes 47 questions above 0.80 overlap and four above 0.95. These shared templates qualify the interpretation of math generalization. The fixed-map code and conversation evaluations give additional task coverage without changing the fitted map.

F FITTING TIME AND ISOLATED MEMORY

Table 25: Selected map size and warm fitting-loop duration on four RTX 6000 Ada GPUs. Models are loaded before the recorded loop begins.

Drafter	Target	Trainable weights (millions)	Fitting loop (seconds)
DFlash	8B	52.43	85.6
DFlash	14B	65.54	118.0
EAGLE-3	8B	52.43	86.7
EAGLE-3	14B	65.54	118.5

The fitting-loop timer covers source and target feature computation and matrix optimization. A complete setup also loads model weights, writes the map, reloads it and initializes generation. Published original training counts in Table 5 provide context for the additional-data requirement. They are not a wall-time or token-budget comparison with fitting a new native drafter.




Figure 13: Isolated EAGLE-3 peak allocated memory on RTX 6000 Ada. Each method runs in a fresh process on eight paired prompts per target. Source removal saves 7.80 GiB at 8B and 7.63 GiB at 14B.

GPU peak counters are reset before the measured request. Allocated memory counts live tensor allocations and differs from reserved allocator memory. The mixed-arm processes used for timing are not the basis of this source-removal measurement.

G ADDITIONAL RELATED WORK

Improving drafting. EAGLE predicts future features with shifted tokens (Li et al., 2024a). EAGLE-3 combines feature layers and simulated drafting steps (Li et al., 2025). DFlash predicts a block in parallel (Chen et al., 2026), while Medusa uses several prediction heads (Cai et al., 2024). GRIFFIN addresses training-to-decoding token alignment (Hu et al., 2025). RepSpec changes the training structure (Huo et al., 2026), while VSD and Draft-OPD optimize accepted proposals (Zou et al., 2026; Lei et al., 2026). RelaySpec inherits the trained drafter and changes only its conditioning input.

Proposal structure and execution. EAGLE-2 and OPT-Tree adapt proposal trees (Li et al., 2024b; Wang et al., 2025). SpecInfer organizes tree verification (Miao et al., 2024). Distributed speculative inference and speculative speculative decoding overlap computation (Timor et al., 2025b; Kumar et al., 2026). Draft & Verify skips target layers (Zhang et al., 2024), LayerSkip trains early exits (Elhoushi et al., 2024), and Lookahead checks candidates without an auxiliary drafter (Fu et al., 2024). RelaySpec retains an external drafter and holds its released proposal structure fixed. We evaluate the cost and accepted progress of changing its feature source.